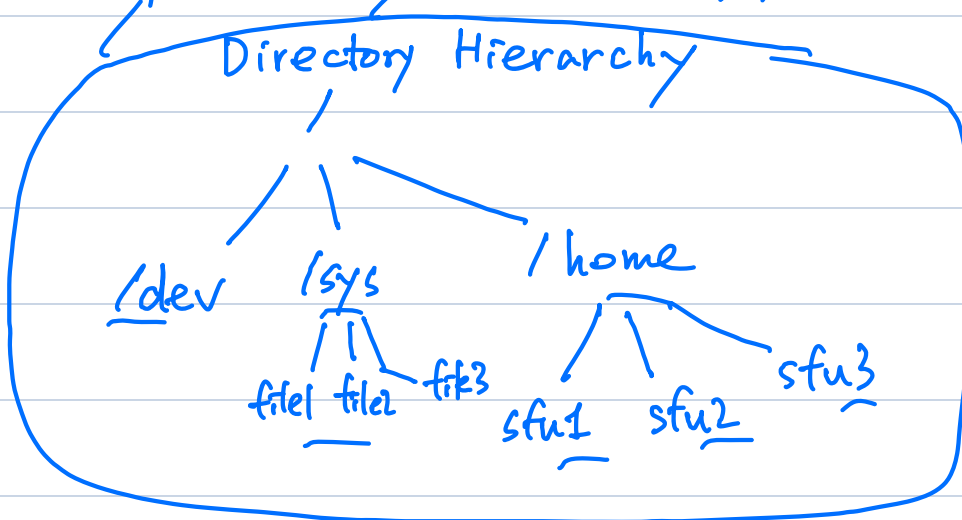


Virtual File System

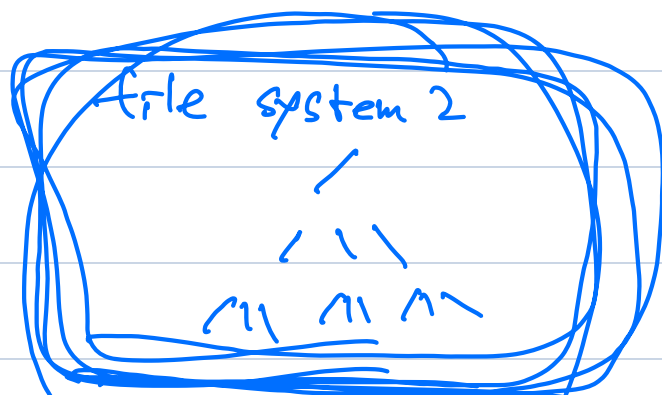
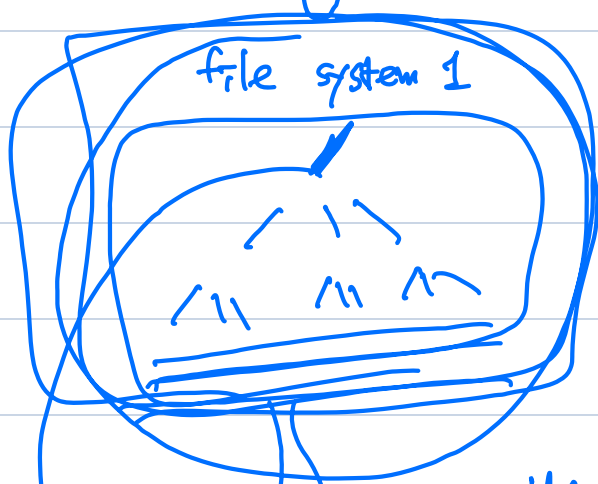
syscall boundary
file abstraction.
syscalls. (open. close.
read. write)



file systems
implement the
interface.



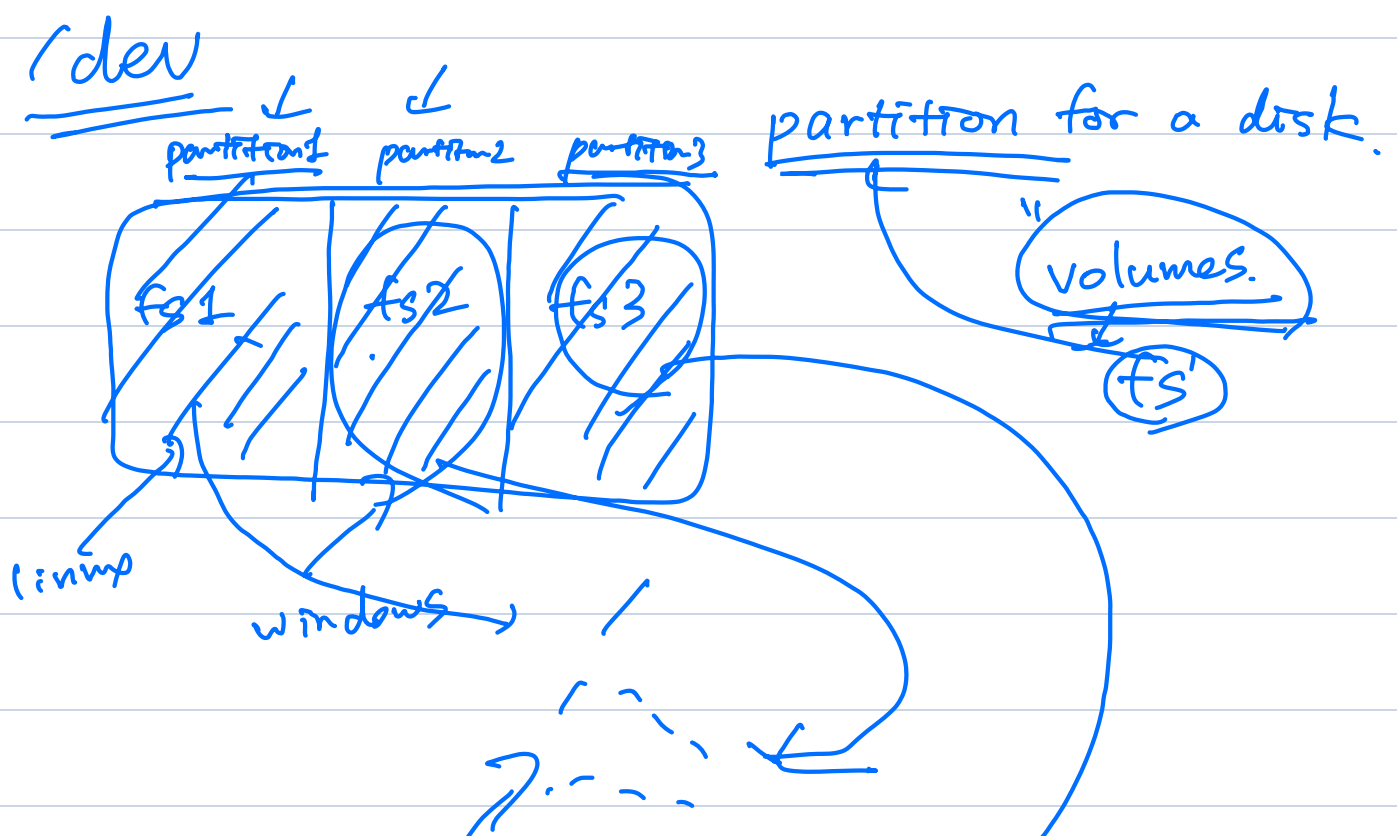
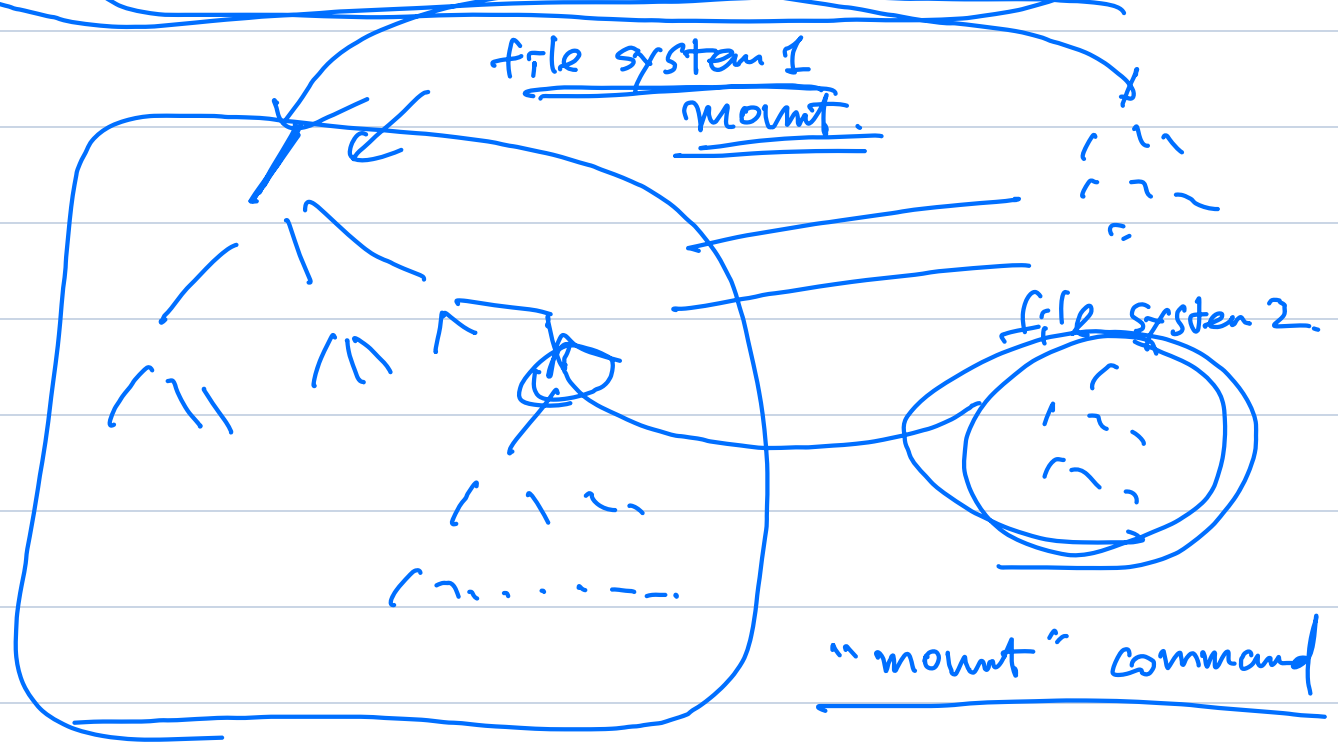
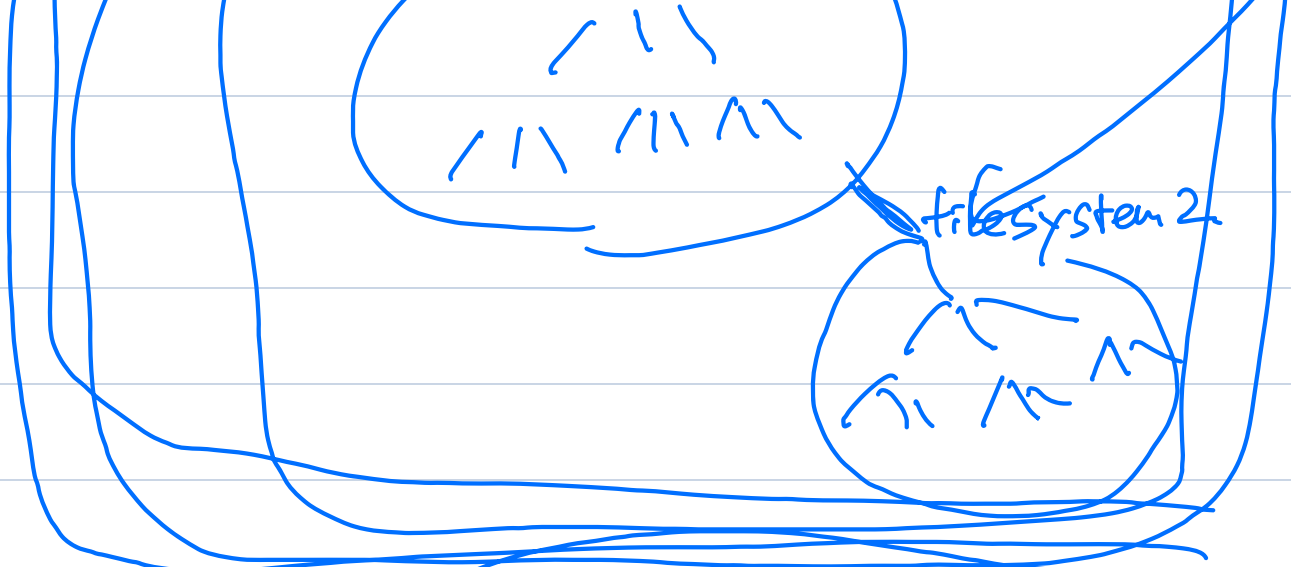
mounting / unmounting



User-facing directory hierarchy

mounting done





MBR

disk block : 4KB. (matches the page size)



* Files: data content

* Directories: considered files.

content is the files that they contain.

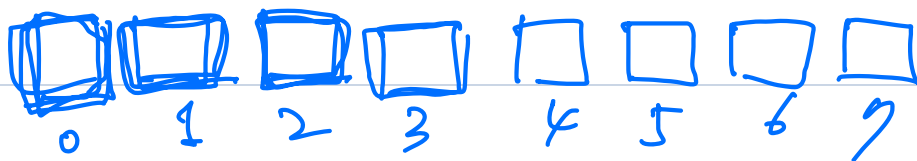
metadata: i-nodes

data: ~~the~~ file contents

data blocks

metadata blocks (i-nodes)

* Example



: meta data.

* i-node block

* i-bitmap, d-bitmap

* super block

: data blocks



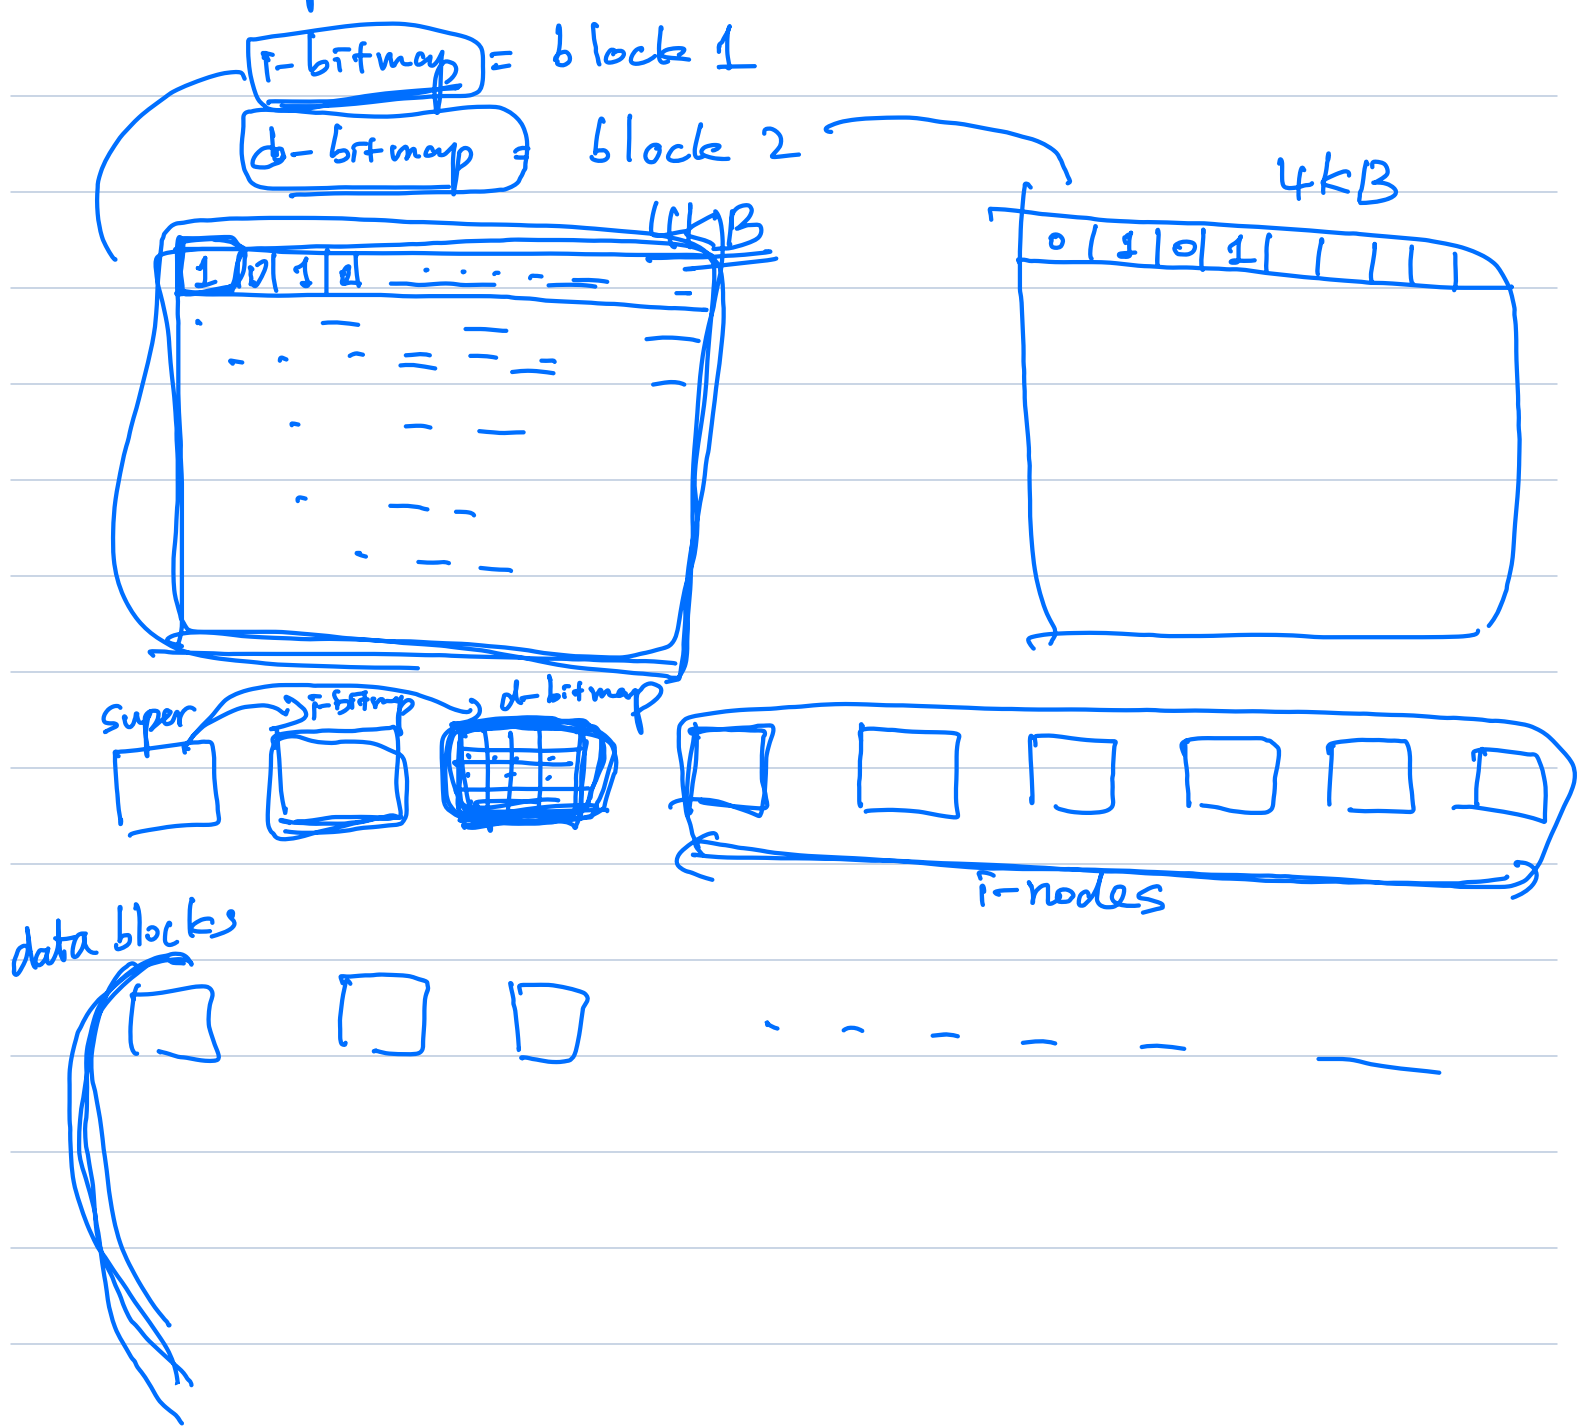
* meta data blocks

* i-node blocks

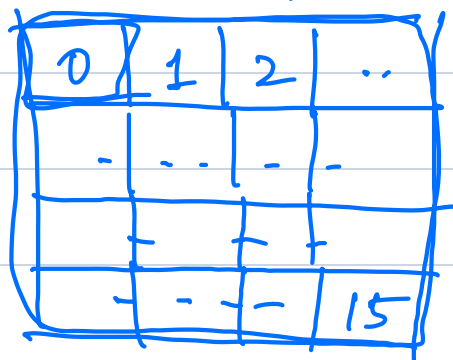
* i-bitmap, d-bitmap

* Superblock (first block)
block 0

* Example.



* i-node block 4k



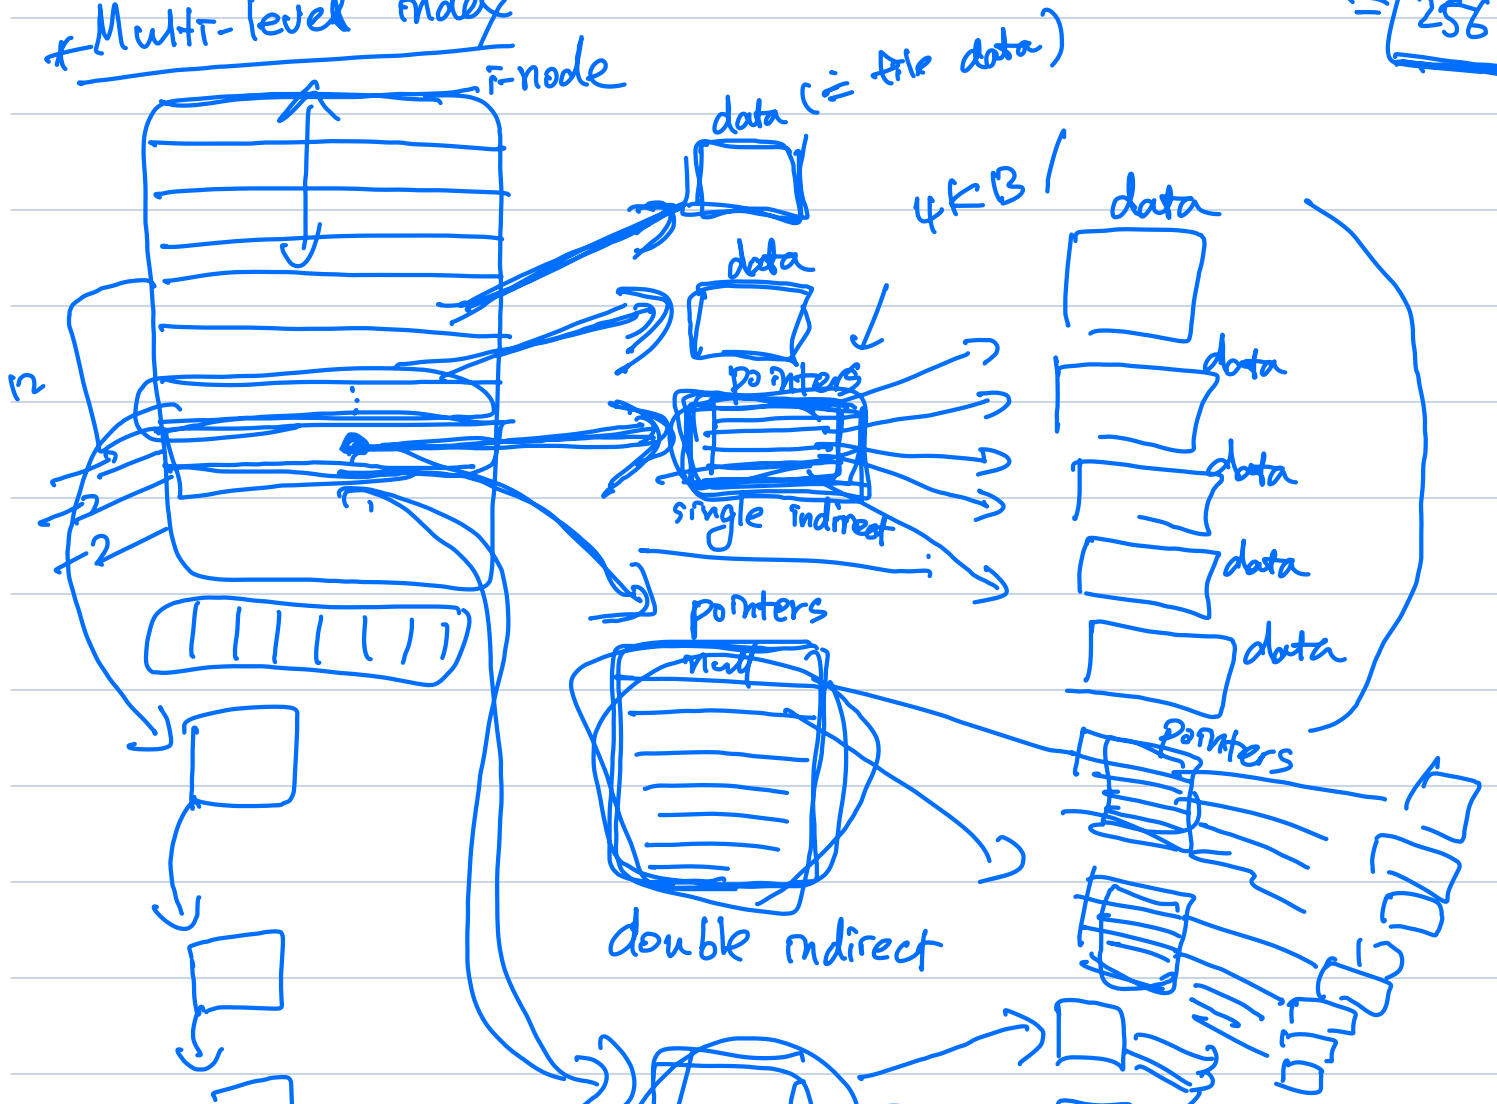
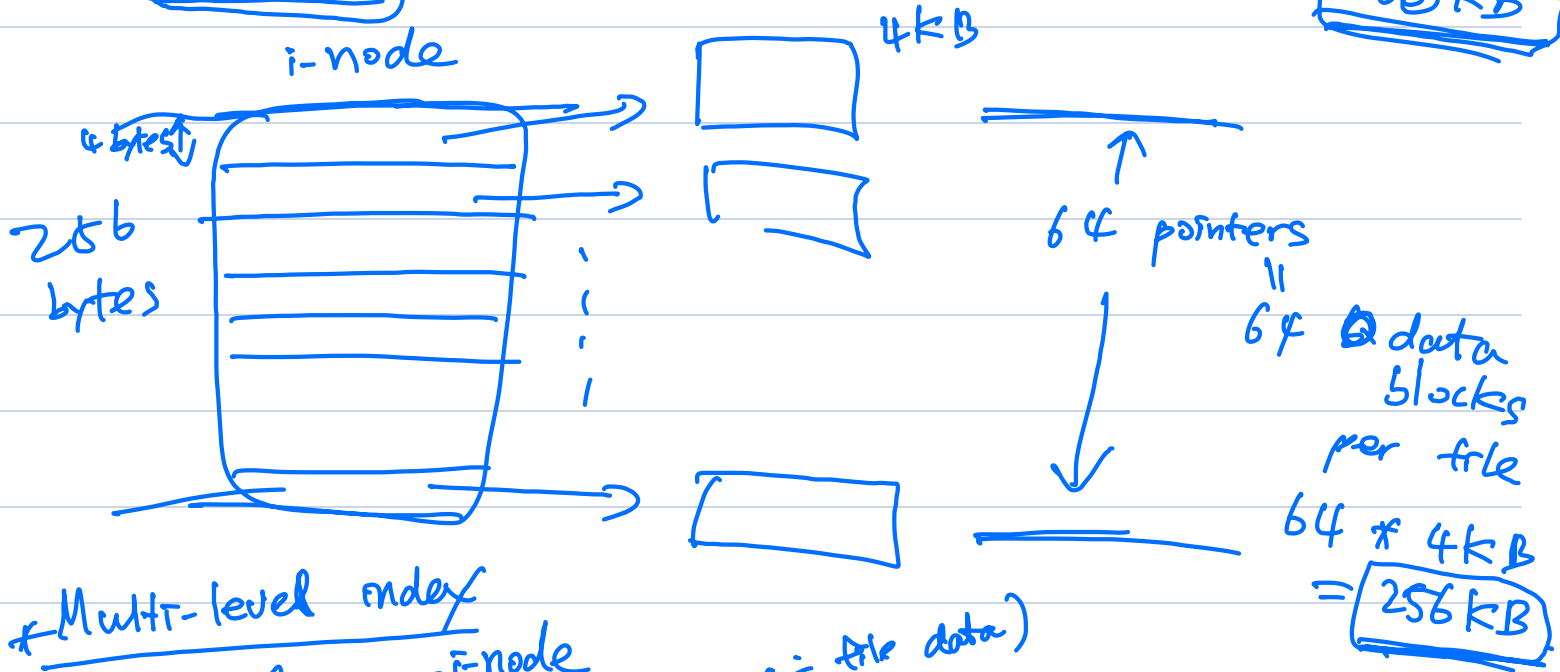
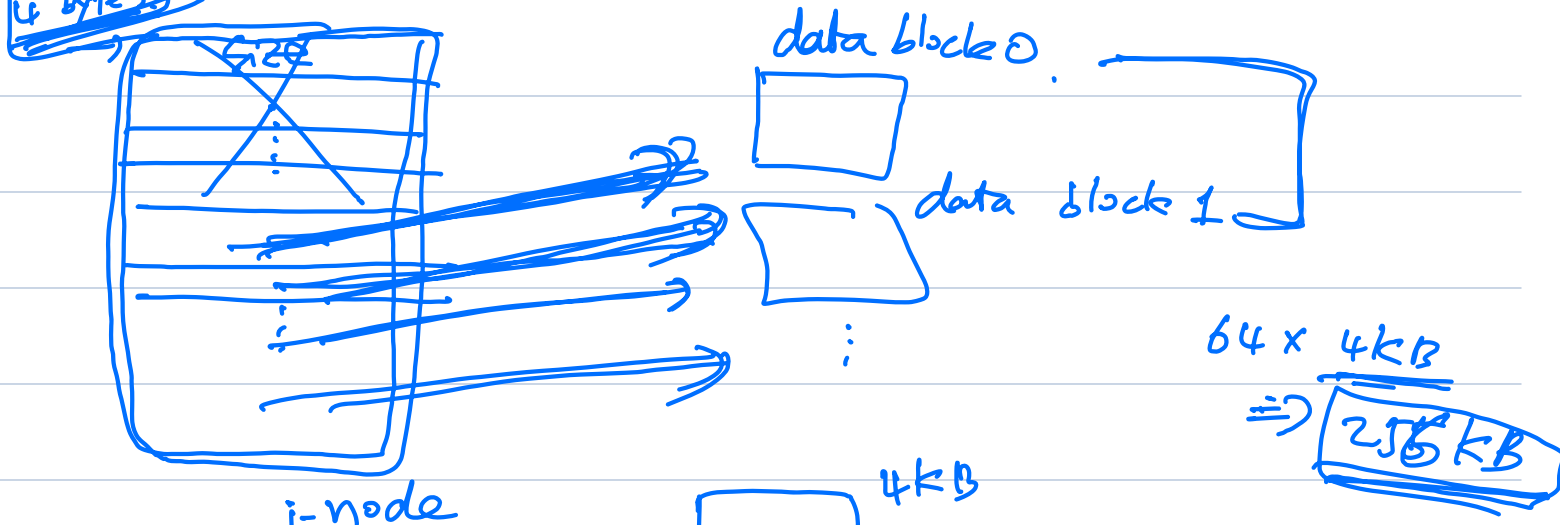
256/4

- * size
- * ~~time~~ last accessed time
- * time created
- * permission
- ...

* file name is not stored.
(stored in the directory file that contains this file)

* where the data blocks are

i-node 256 bytes



* Q: Assume block size: 4KB

pointer size: 4 bytes

- 12 direct entries (pointers)
- 1 single indirect entry (pointer)
- 1 double indirect entry (pointer)
- 1 triple indirect entry

what's the maximum file size?

A: 12 * 4KB

4KB block / 4 = 1K.

1024 * 4KB

1024 pointers

1024 * 1024 * 4KB

4 bytes

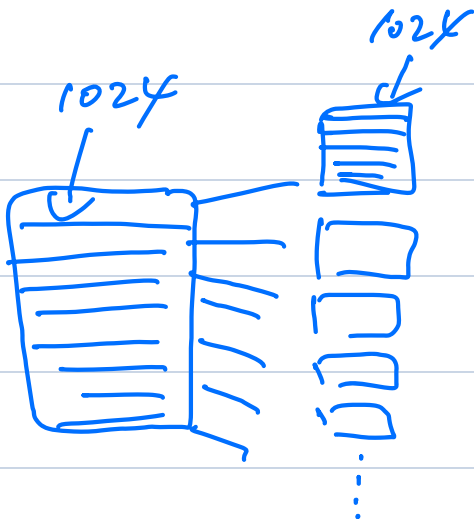
1024 * 1024 * 1024 * 4KB

" 4TB

4KB



4KB / 4B
= 1K
= 1024



= 1024 * 1024

open:

* /foo/bar

i-node #2

read root i-node

read root data block →

file 0 → i-node

file 1 → i-node

foo → i-node

read foo i-node

read foo data block →

(file2 → i-node
file3 → i-node
bar → i-node
⋮)

read bar i-node

fd = open ("/foo/bar");

↘ bar i-node

read (fd, buf, ...)

write (fd, buf, ...)

* read (fd)

read bar i-node

read bar data

write bar i-node

* write (fd)

read bar i-node

read data bitmap

write data bitmap

write data block

write bar i-node

* create "/foo/bar"

• read root i-node

read root data

read foo i-node

read foo data

(foo → i-node)

(~~bar~~)

— read r-node bitmap
— write r-node bitmap ↙

— write foo data (bar → r-node)

— read bar r-node

— write bar r-node (initialization)

— write foo r-node

* caching. (write-back. ~~tr~~
write-through
