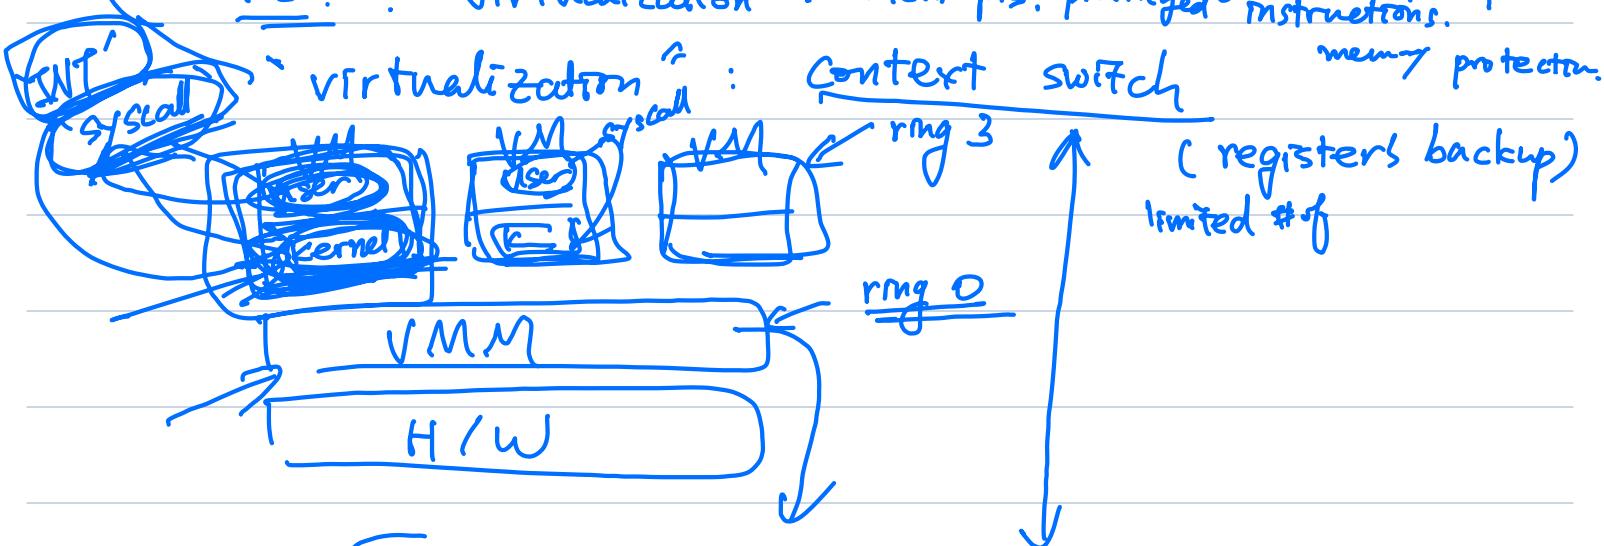


CPU : privileged instructions that user-level programs can't run

Memory : "Virtualization": Virtual memory directly.

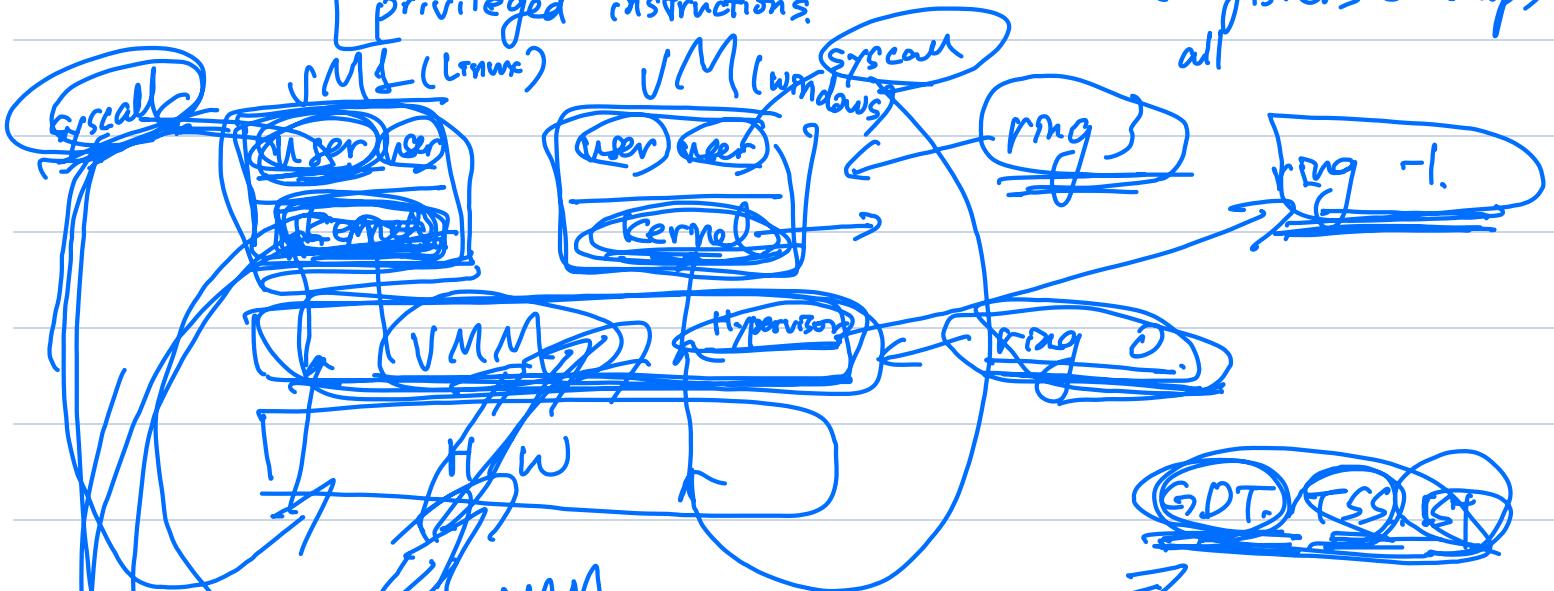
I/O : "virtualization": kernel-reserved memory (page table, per-process)



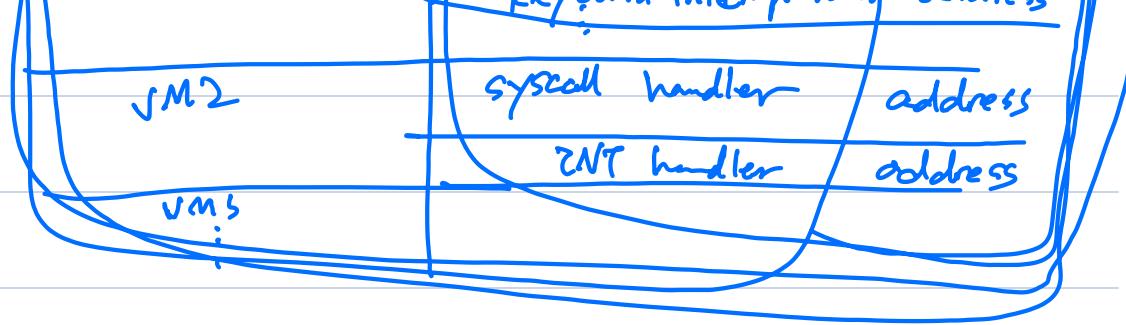
CPU : "virtualization" : machine switch

privileged instructions.

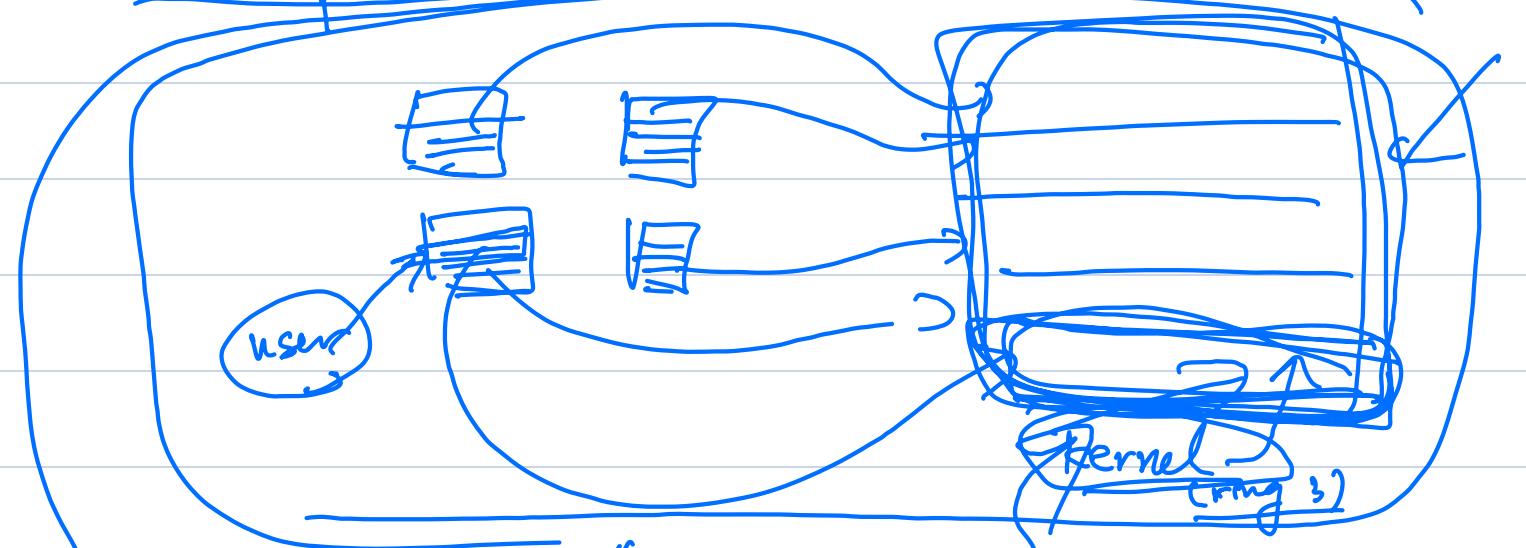
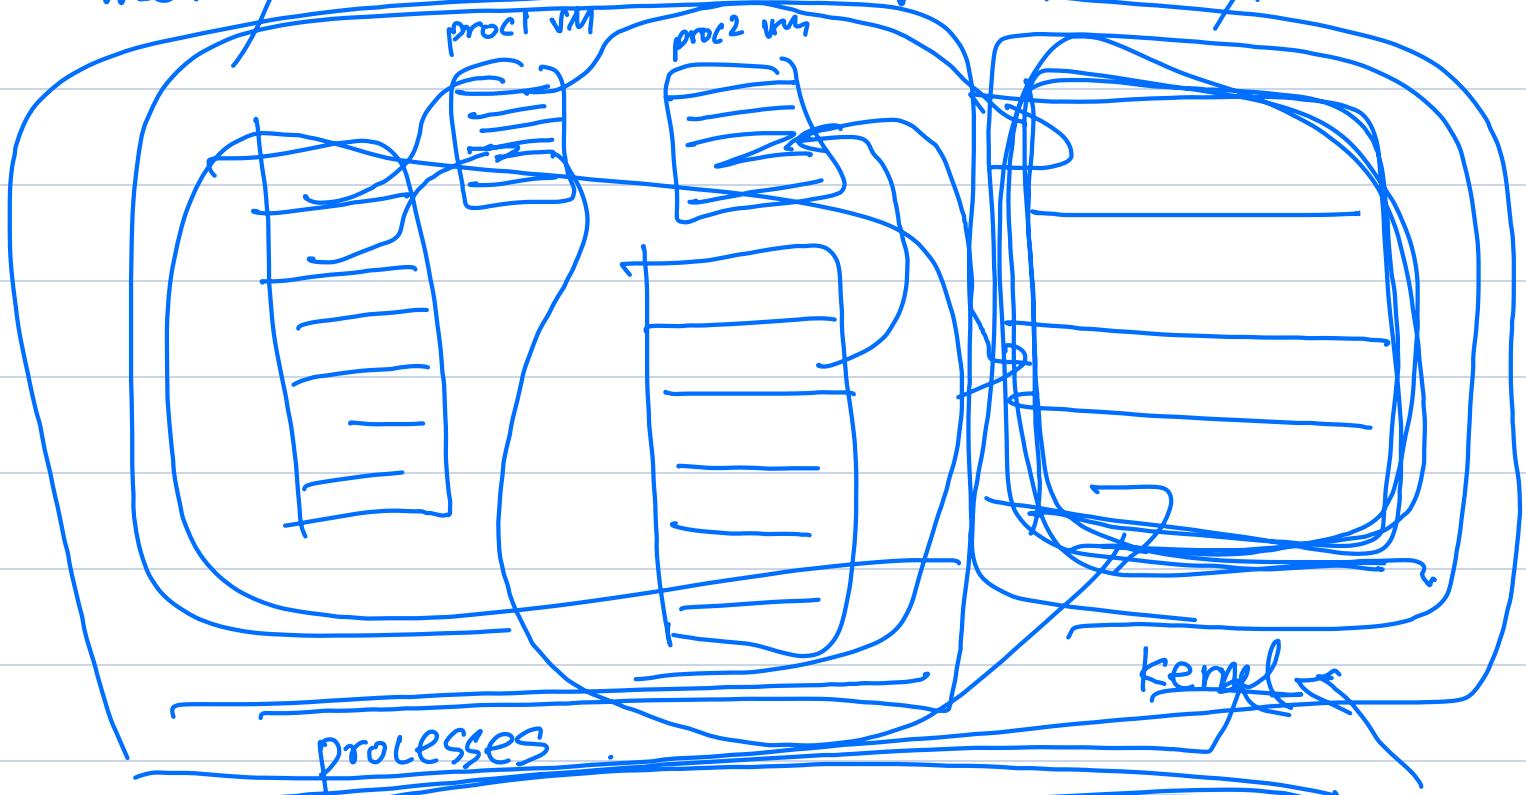
(registers backup)
all



syscall handler address	IA32 Handler address	address
IA32 Handler address	newboard interrupt handler address	



memory : virtualization . virtual memory.

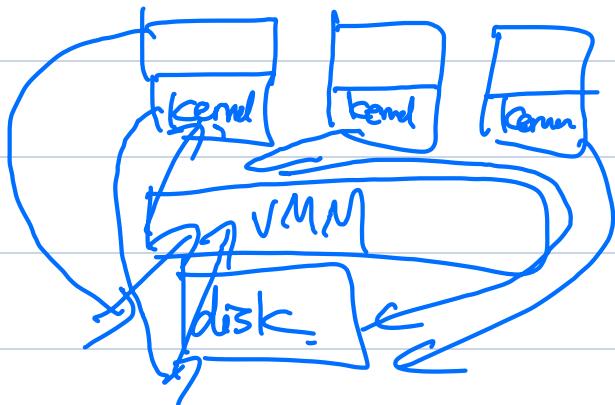




I/O

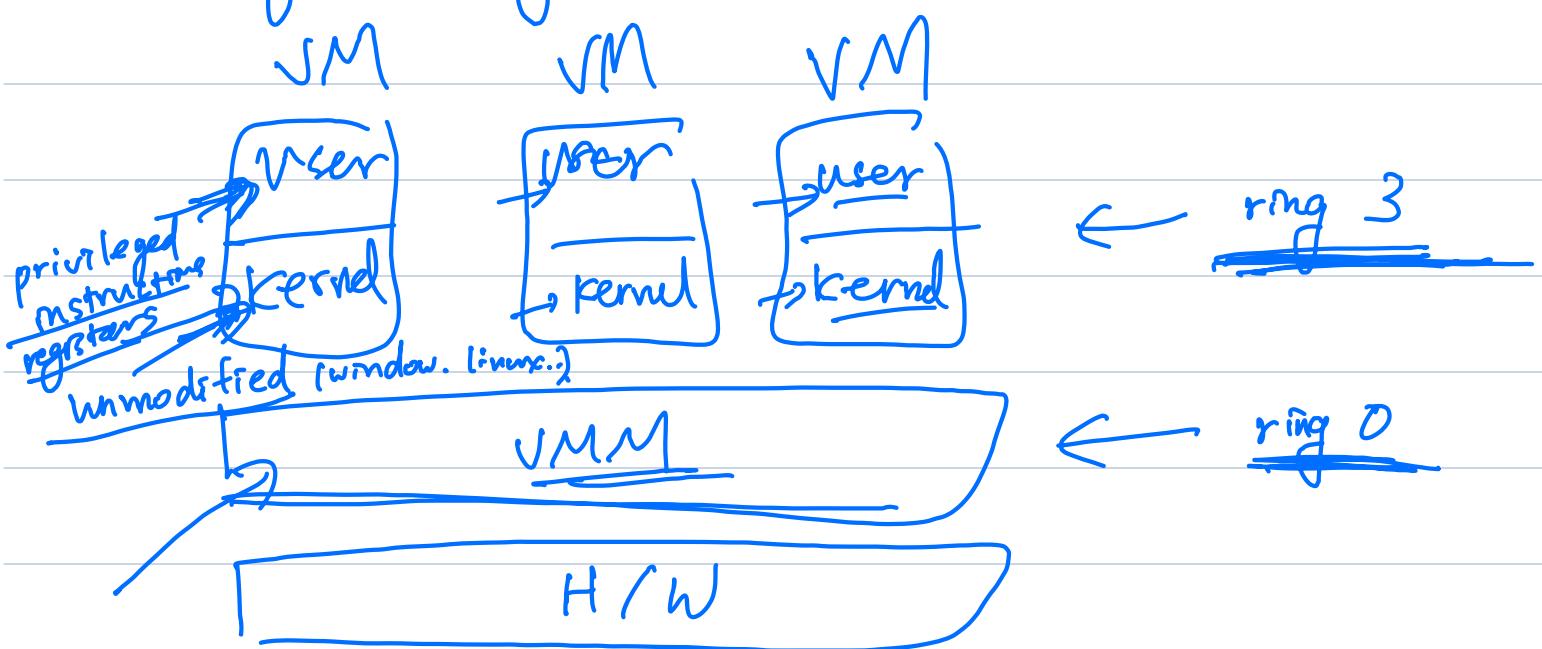
* - I/O instruction (User program)

- trap.
- VMM finds a handler within the corresponding kernel.
- kernel handles it..



IN port & data

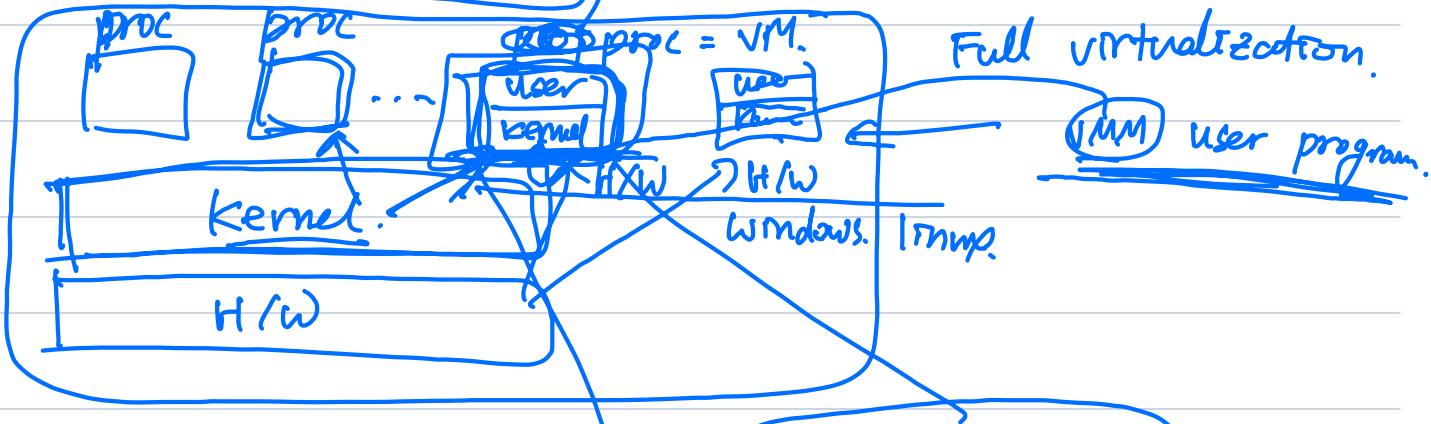
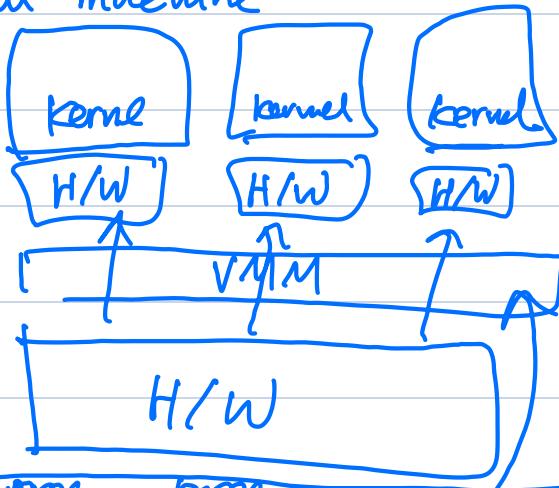
* ring 0 vs. ring 3.



- * problem with ring 3
 - * Kernel memory (pages) now user accessible.
 - * Performance: kernel can't have privileged access \Rightarrow VMM needs to handle traps.
 - * Instructions that behave differently based on the privilege level.
`popf` \Rightarrow safely fails in ring 3.
in ring 0. ok

* Binary translation

↑ Full machine virtualization



CPU:

mem.

T10

User programs.

Kernel

VMM

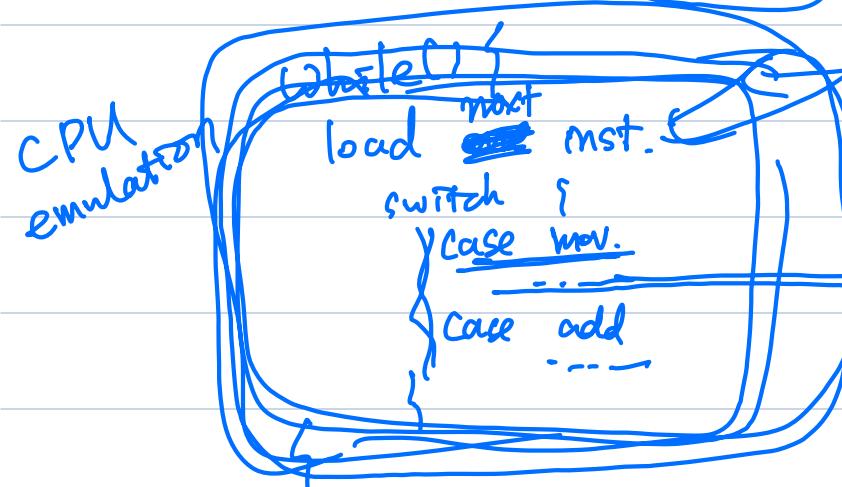
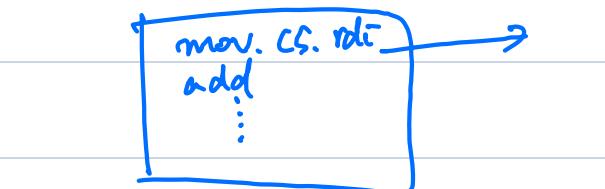


10

PS2 emulator

Intercept instructions.

* Emulate the access
to CS. SS.



CPU state
(software)

ub4 CS;
ub4 SS.;
ub4 rdi

CS = rdi

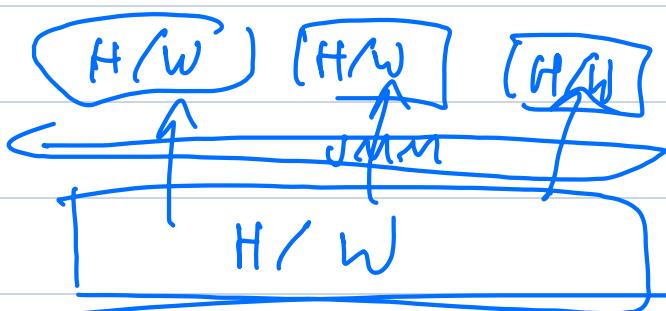
emulator

- Java VM

- Python VM.

⇒ Java
bytecode.

* Containerization



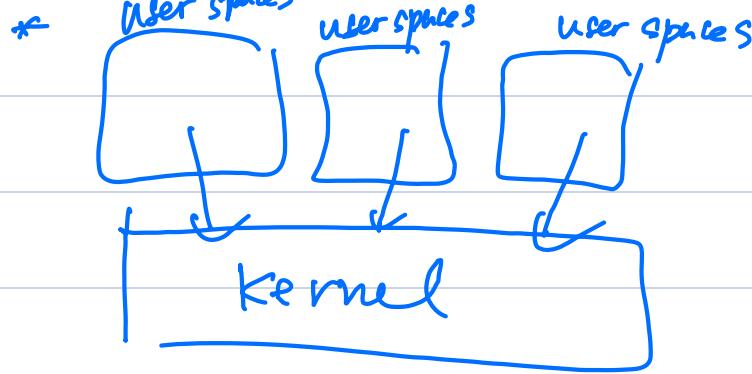
User programs

kernel

H/W

Containerization.
(kernel virtualization)
or user space virtualization

JMM virtualization



* User space -

* processes. (UI-shell, other user programs, ...)

+ File system.

+ I/O -

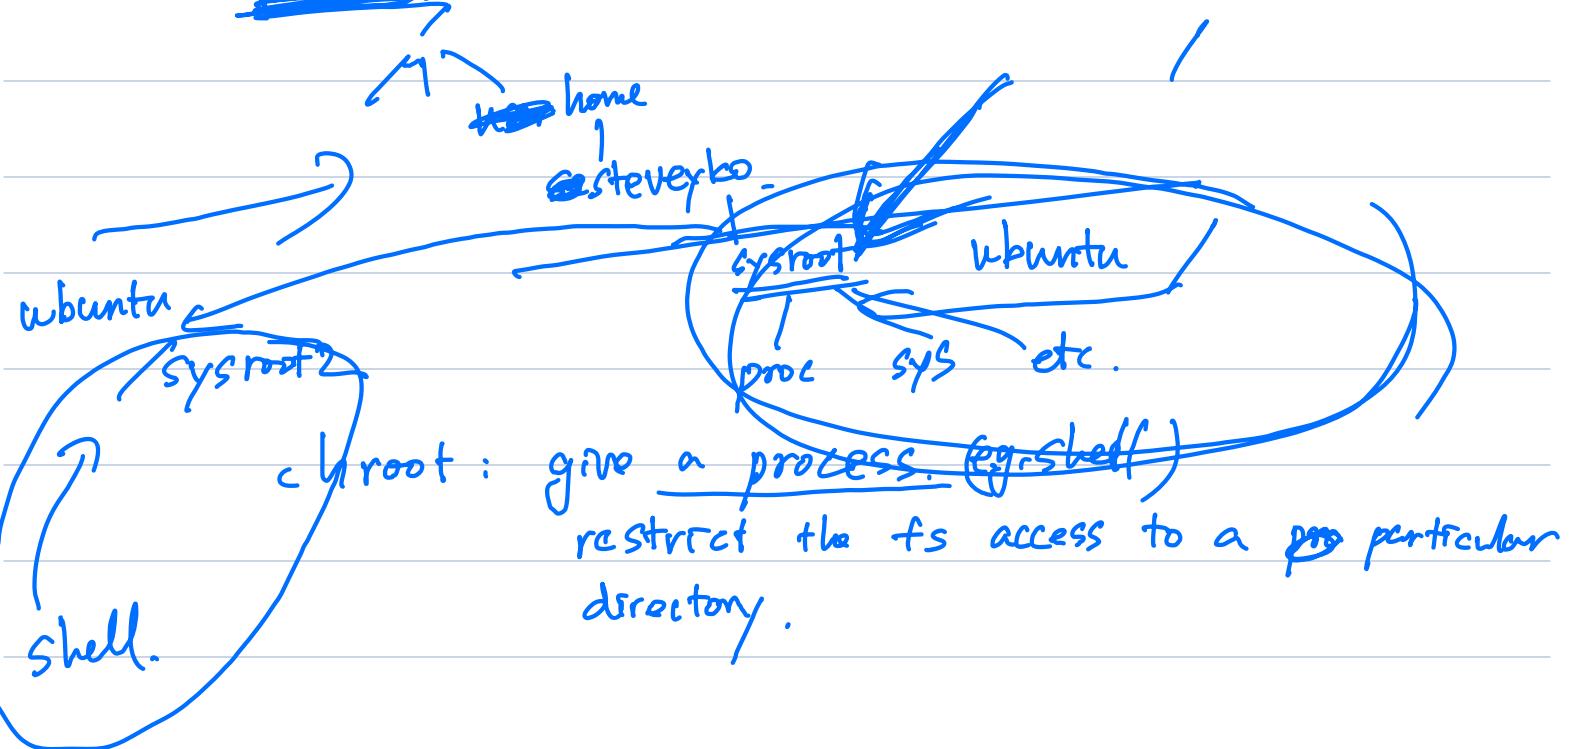
* chroot, namespace, cgroups.

↓
file system

↓
processes

↓
I/O

* chroot.



* Name space -

* restrict a process's view of other processes.

* Eop. ps.

